

## Lecture 12 (Transport 2)

# TCP Implementation

CS 168, Spring 2025 @ UC Berkeley

Slides credit: Sylvia Ratnasamy, Rob Shakir, Peyrin Kao

# Implementing TCP: Byte Notation (Segments, Sequence Numbers)

---

Lecture 12, CS 168, Spring 2025

## Implementing TCP

- **Byte Notation  
(Segments, Sequence Numbers)**
- Maintaining State (Full Duplex,  
Connection Setup and Teardown)
- Sliding Window
- Header

## Notation Change: Bytes vs. Packets

---

So far, we've used *packets* as the primary unit of data.

- Each packet has a number.
- Acks reference packet numbers.
- Window size expressed in terms of number of packets.

TCP is implemented with *bytes* as the primary unit of data.

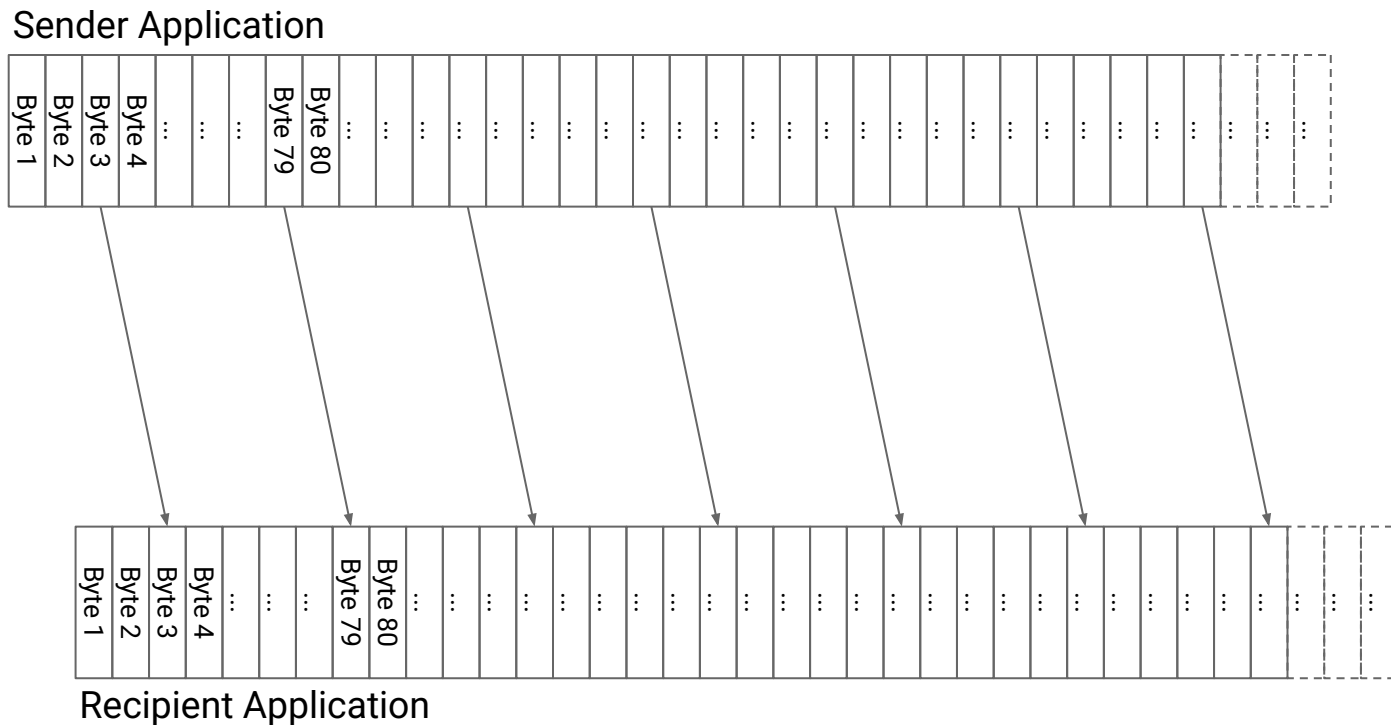
- Each byte has a number.  
Packets are defined by the number of the first byte inside.
- ACKs reference byte numbers.
- Window size expressed in terms of number of bytes.

You should be prepared to reason in terms of either.

# TCP Segments

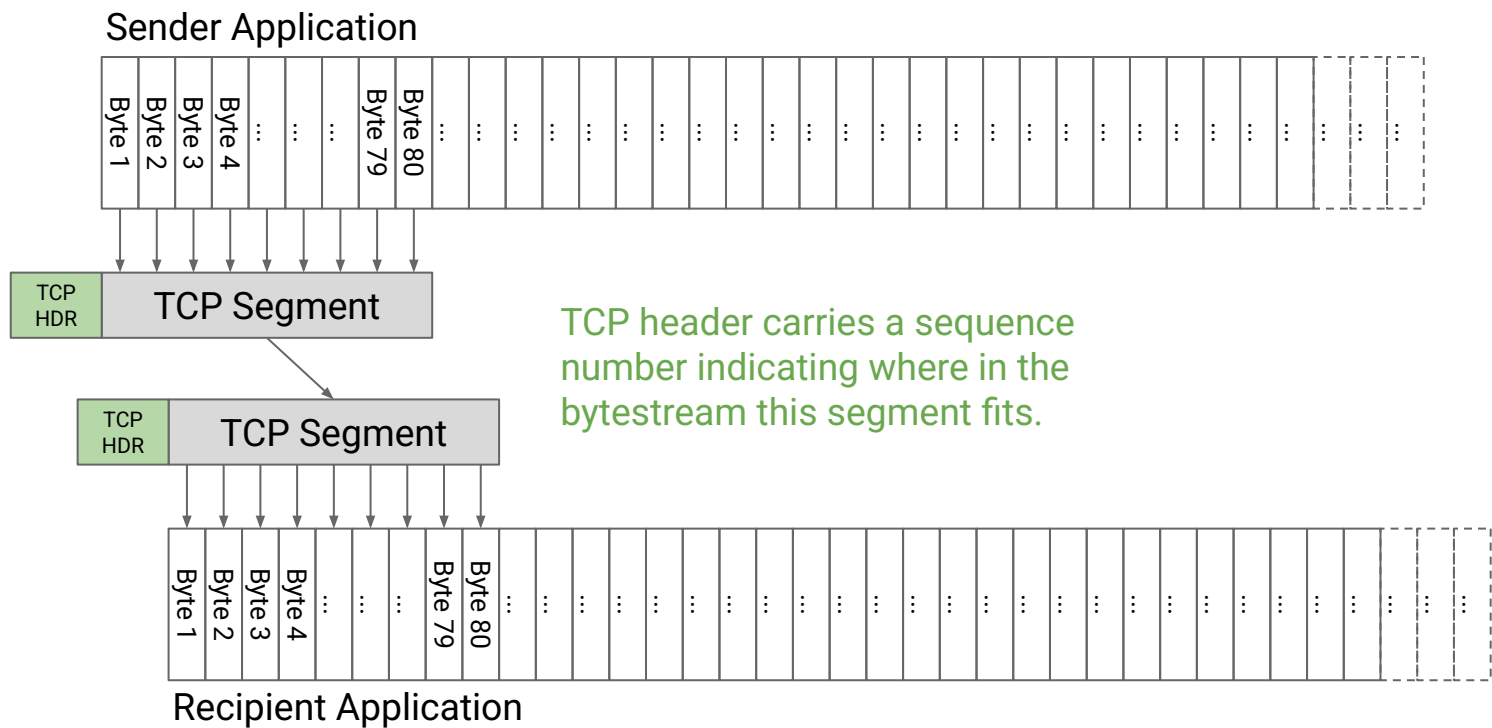
TCP provides a reliable, in-order bytestream.

We have to split this bytestream into packets.



# TCP Segments

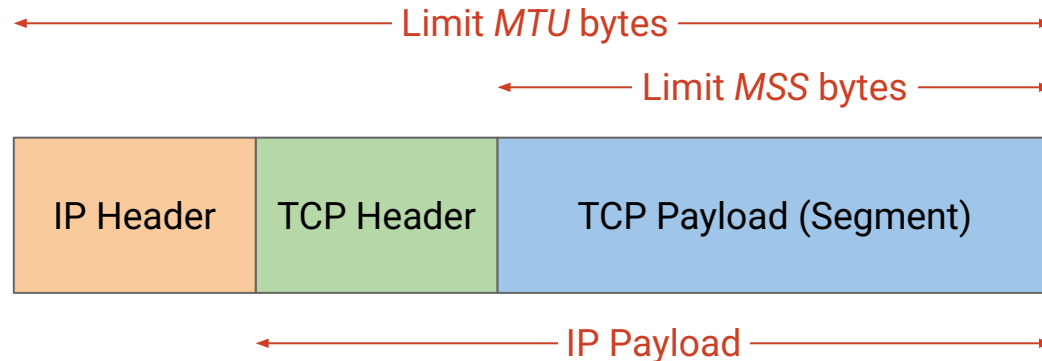
A segment is sent when the segment is full (max segment size),  
or when the segment is not full, but times out waiting for more data.



**TCP/IP packet:** IP packet with TCP header and TCP data inside.

Size limits:


- IP packet: Maximum transmission unit (MTU).
- TCP segment: Maximum segment size (MSS).
- $MSS = MTU - (IP\ header) - (TCP\ header)$ .




## TCP Sequence Numbers

Numbering starts at a randomly-generated **Initial Sequence Number** (ISN).

- First byte is  $ISN+1$ , then  $ISN+2$ , etc.
- Starting at a randomly-chosen ISN is very important for security!



Sender Application

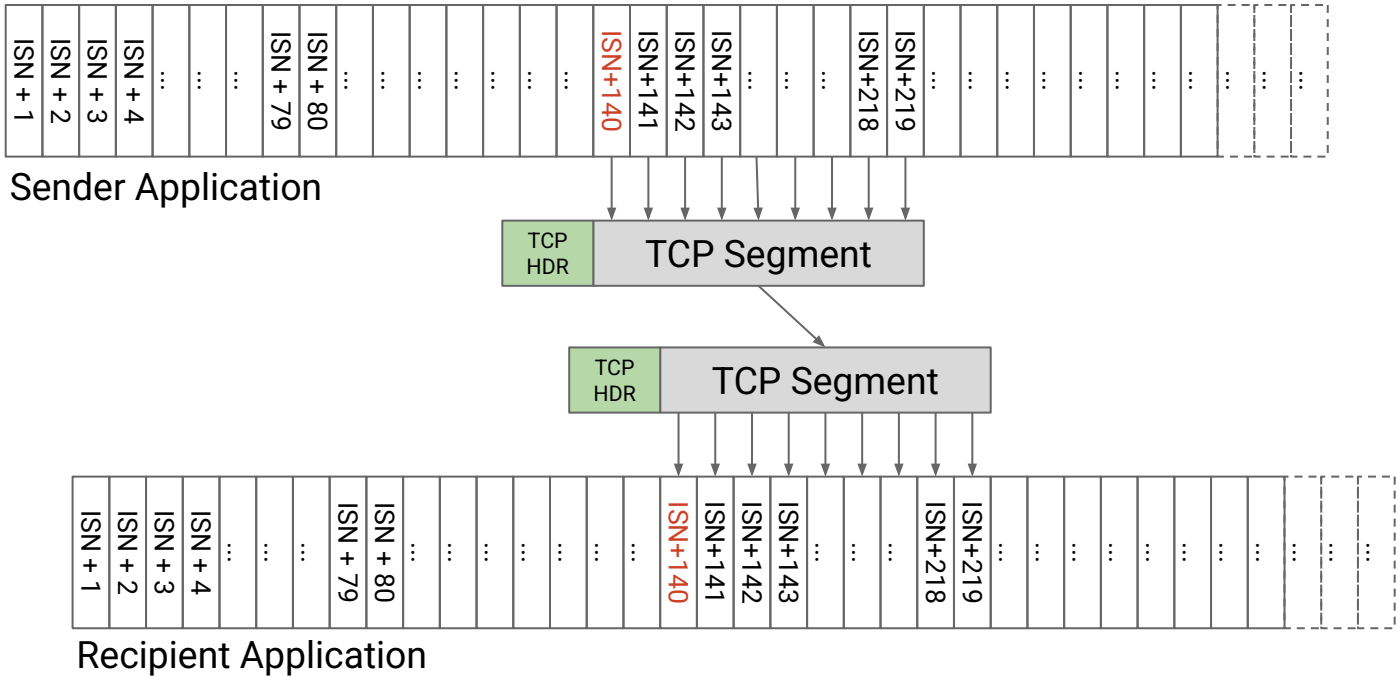


Recipient Application

# TCP Sequence Numbers

The sequence number of a segment is the number of the *first byte* in the segment.

Example: In the segment below, the sequence number is *ISN + 140*.

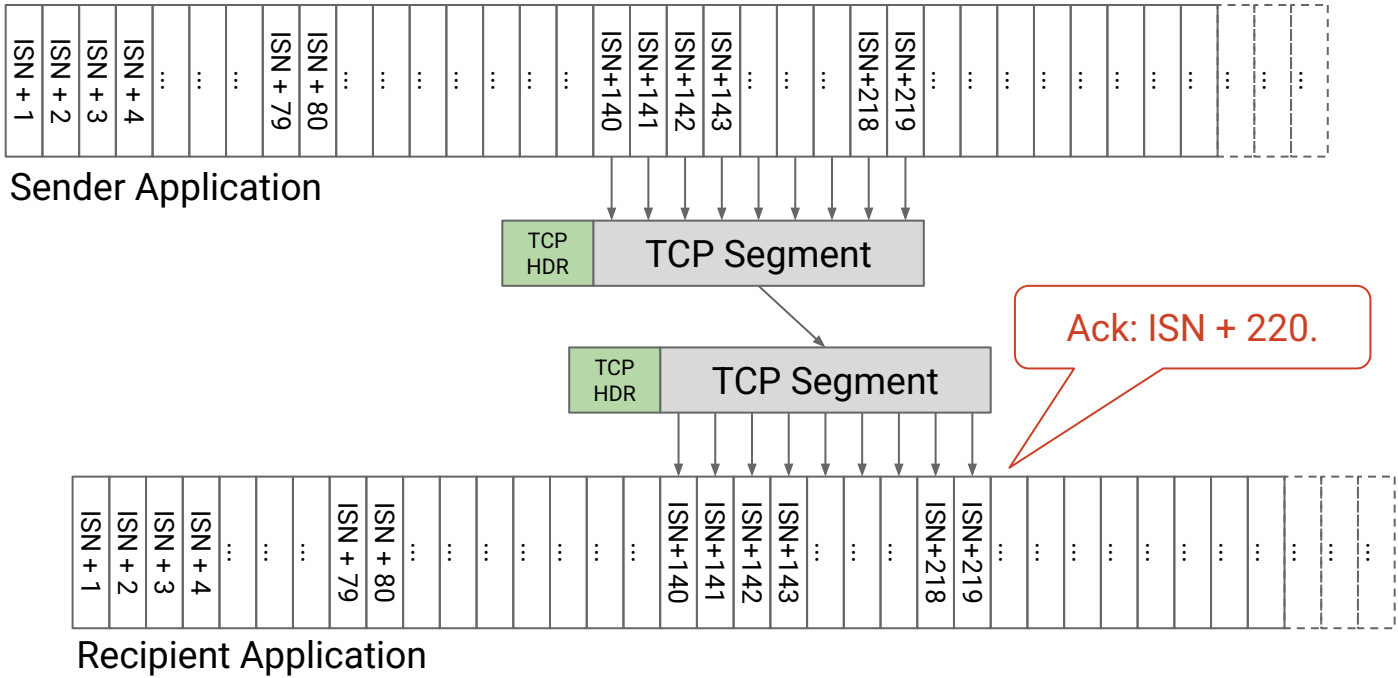




# TCP Ack Numbers

The ack number indicates the next expected byte (i.e. the first unreceived byte).

Example: All bytes up to (and including) *ISN* + 219 have been received, so the next unreceived byte is *ISN* + 220.



## TCP Sequence and Ack Numbers

---

The sequence number of a segment is the number of the *first byte* in the segment.

- Sender sends a packet with sequence number  $j$ .
- The packet contains  $B$  bytes.
- Bytes in the packet are numbered:  $j, j+1, j+2, j+3, \dots, j+B-1$ .

Recipient sends a cumulative ack (number of highest byte received, plus one).


- If all prior data is received, ack number is  $j+B$ .
- Think of this as the next expected byte, or the first unreceived byte.
- If earlier data before this packet is missing, the ack number will be lower.

## TCP Sequence and Ack Numbers

Assuming only one packet in flight, all packets length  $B$ , and no loss:

The last ack number is equal to the next sequence number.

- "I expect  $j + B$  next."  $\rightarrow$  "I'm sending  $j + B$ ."



# Maintaining State (Full Duplex)

---

Lecture 12, CS 168, Spring 2025

## Implementing TCP

- Byte Notation  
(Segments, Sequence Numbers)
- **Maintaining State (Full Duplex,  
Connection Setup and Teardown)**
- Sliding Window
- Header

Reliability requires maintaining *state* at the end hosts.

- Sender has to remember:
  - Which packets have been sent and not acked?
  - How much longer on the timer before I resend a packet?
- Receiver has to buffer the out-of-order packets.
- State is maintained at the end hosts, not in the network.

In each separate connection, both end hosts need to maintain state.

# TCP is Full-Duplex

So far, we defined a sender and a recipient in every connection.

Connections in TCP are **full-duplex**.

- Both hosts can send data, and both hosts can receive data.
- A can send to B, and B can send to A, simultaneously, in the same connection.

To support full-duplex connections:

- Two sets of sequence numbers: One for A→B bytes, and one for B→A bytes.
- Each packet carries both data and ack information.
  - "Here are some bytes starting at 15..."
  - "...Also, I ack receiving all your bytes up to (not including) 84."

15	16	17	18	19	20	21
H	e	l	l	o	,	B

A to B bytestream

77	78	79	80	81	82	83
H	e	l	l	o	,	A

B to A bytestream


## TCP Connection Setup: Three-Way Handshake

---

Goal: Each host tells its ISN to the other host.

1. **SYN**. Client says: "Here's my ISN."
2. **SYN-ACK**. Server says: "I received your ISN. Also, here's my ISN."
3. **ACK**. Client says: "I received your ISN."

After the three-way handshake, both sides can start sending data.




## TCP Connection Setup: Three-Way Handshake

Goal: Each host tells its ISN to the other host.

1. **SYN**. Client says: "My ISN is  $x$ ."
2. **SYN-ACK**. Server says: "I received  $x$  (expecting  $x+1$  next). Also, my ISN is  $y$ ."
3. **ACK**. Client says: "I received  $y$  (expecting  $y+1$  next)."

After the three-way handshake, both sides can start sending data.






## TCP Connection Teardown

Normal termination:


- Each side sends a FIN packet to say: "I'm done sending, but will keep receiving."
- FIN packets must be acked, just like any other data.
- When only one side has sent FIN, the connection is *half-closed*.
- When both sides have sent FIN (both done sending), the connection is closed.



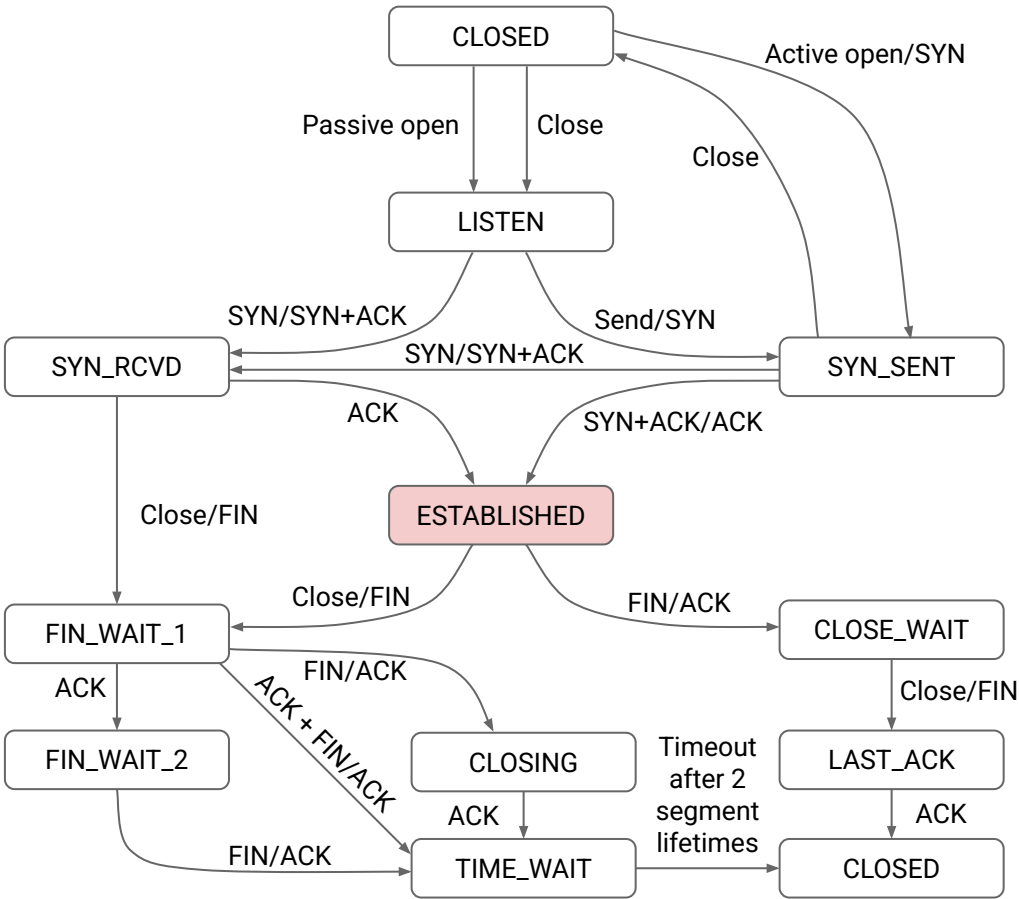
## TCP Connection Teardown

Abrupt termination can be used instead (e.g. in case of error).

- Send a RST (reset) to say: "I will no longer send or receive data."
- RST packets do not need to be acked.
- Any data in flight is lost.
- If the RST sender receives more data later, send another RST.

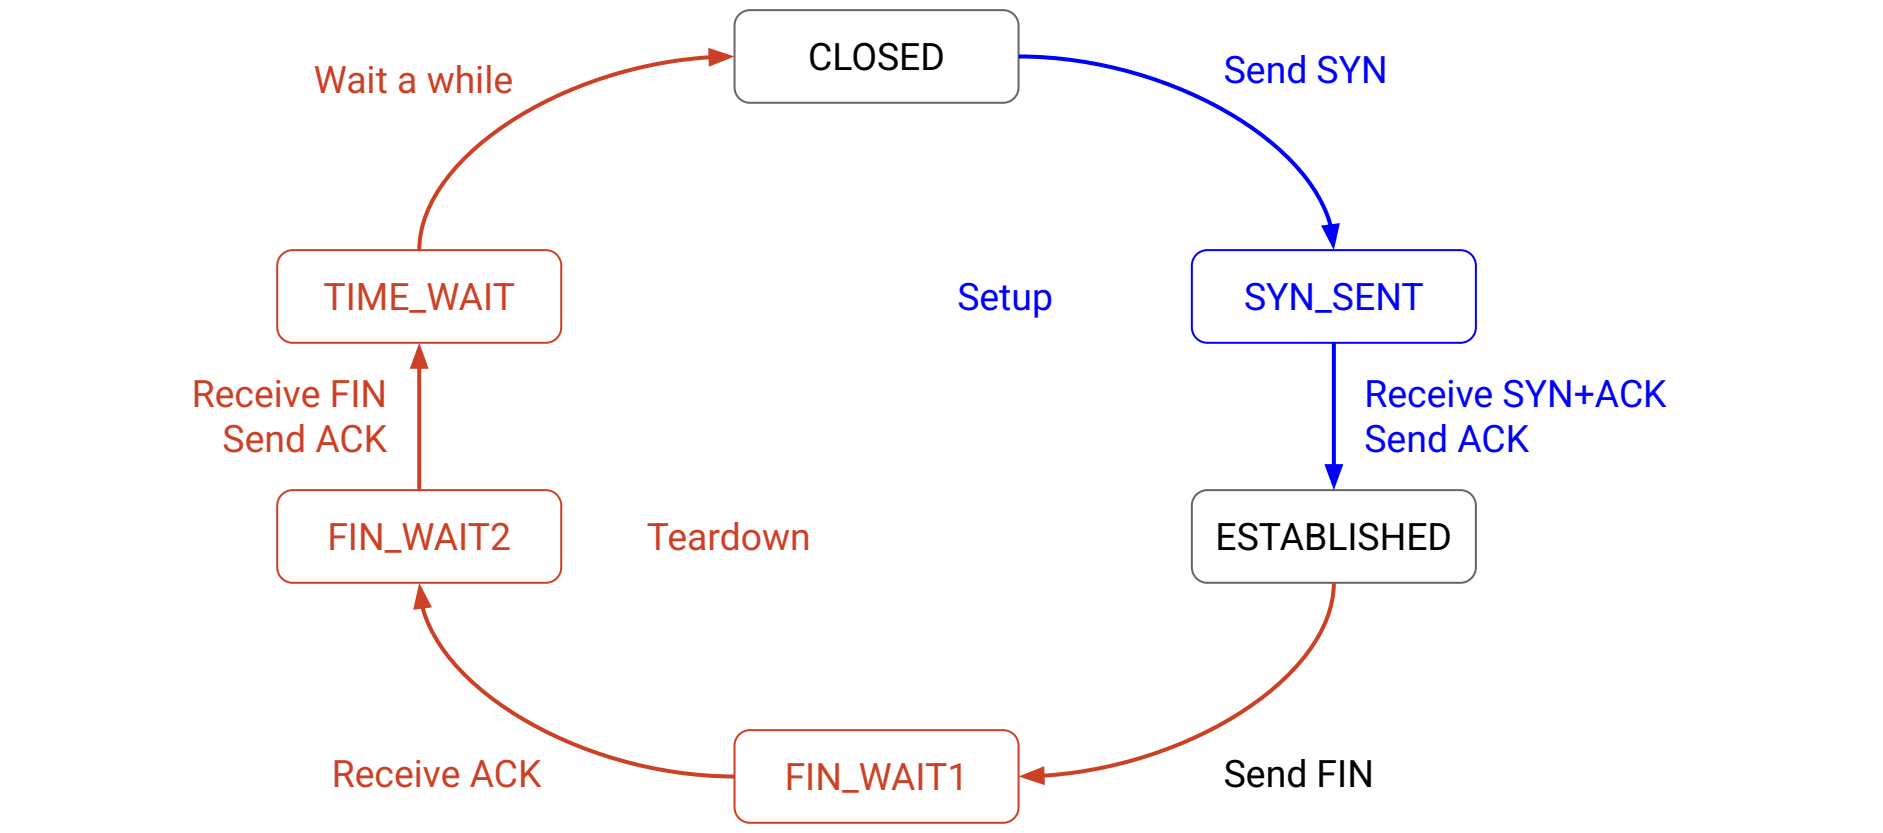


# TCP Setup/Teardown Transition Diagram



The **ESTABLISHED** state is where all data is sent and acked.

# TCP Setup/Teardown Transition Diagram (Simplified)



With full-duplex, if we get a packet but have no data to send, we have two choices:

- Send the ack, with no data.
- **Piggybacking**: Wait for some data, and send the ack with the data.

Piggybacking can be tricky because TCP is in the OS, separate from the application.

- OS doesn't know when application will have more data.
- Application isn't thinking about packets and acks.

SYN-ACKs are always piggybacked.

- Ack and initial sequence number are sent together.
- Not tricky, because OS is doing the handshake, not the application.

# Sliding Window

---

Lecture 12, CS 168, Spring 2025

## Implementing TCP

- Byte Notation  
(Segments, Sequence Numbers)
- Maintaining State (Full Duplex,  
Connection Setup and Teardown)
- **Sliding Window**
- Header

## TCP Sliding Window

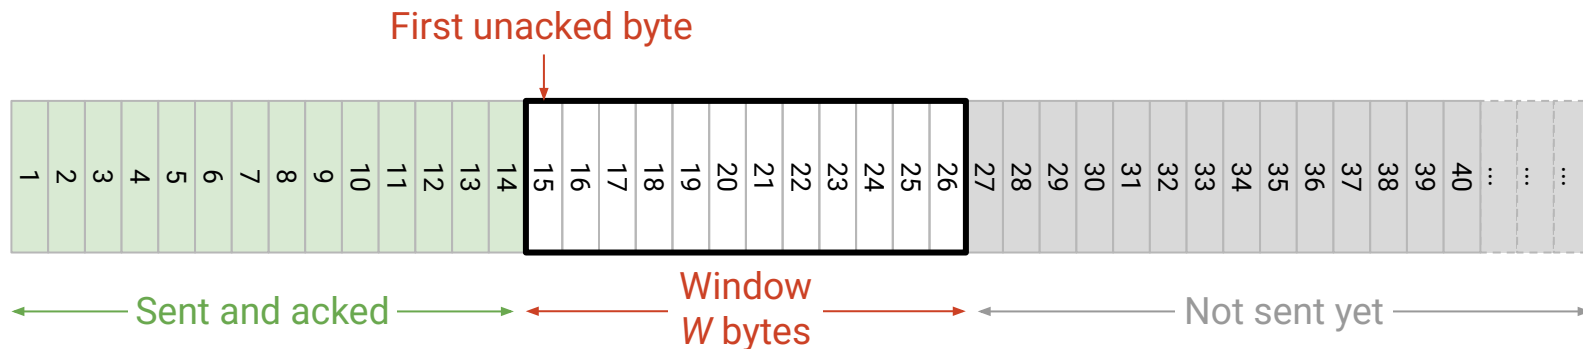
When we measured in packets,  $W$  was the maximum number of packets in flight.

When we measure in bytes,  $W$  is the maximum number of *contiguous* bytes in flight.

- The window is a range of  $W$  contiguous bytes, starting at the first unacked byte.
- Only these  $W$  bytes are allowed to be in flight.

The window slides right if and only if its *leftmost* bytes are acked.

- Example: When 15–18 arrive, we can send 27–30.



## TCP Sliding Window

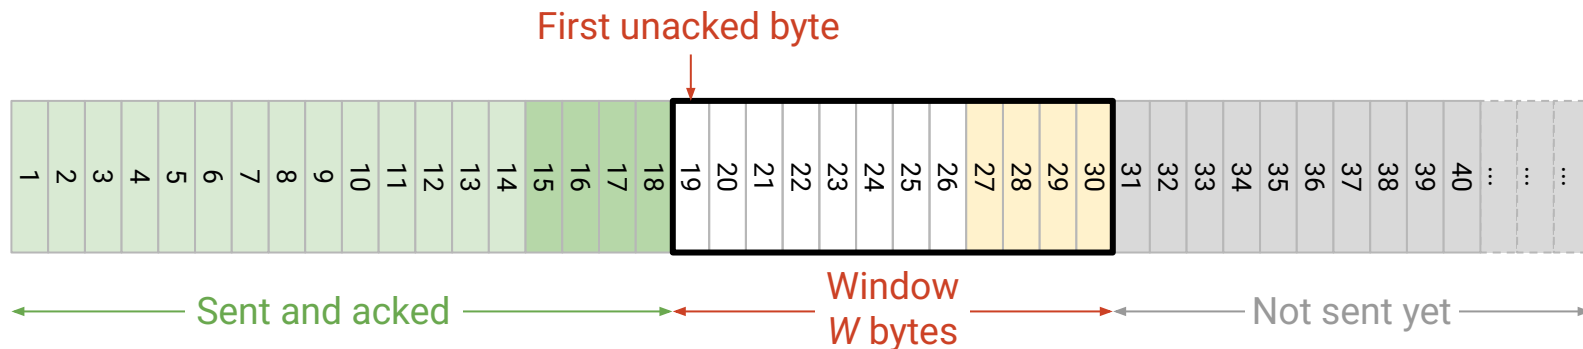
When we measured in packets,  $W$  was the maximum number of packets in flight.

When we measure in bytes,  $W$  is the maximum number of *contiguous* bytes in flight.

- The window is a range of  $W$  contiguous bytes, starting at the first unacked byte.
- Only these  $W$  bytes are allowed to be in flight.

The window slides right if and only if its *leftmost* bytes are acked.

- Example: When 15–18 arrive, we can send 27–30.





## TCP Sliding Window

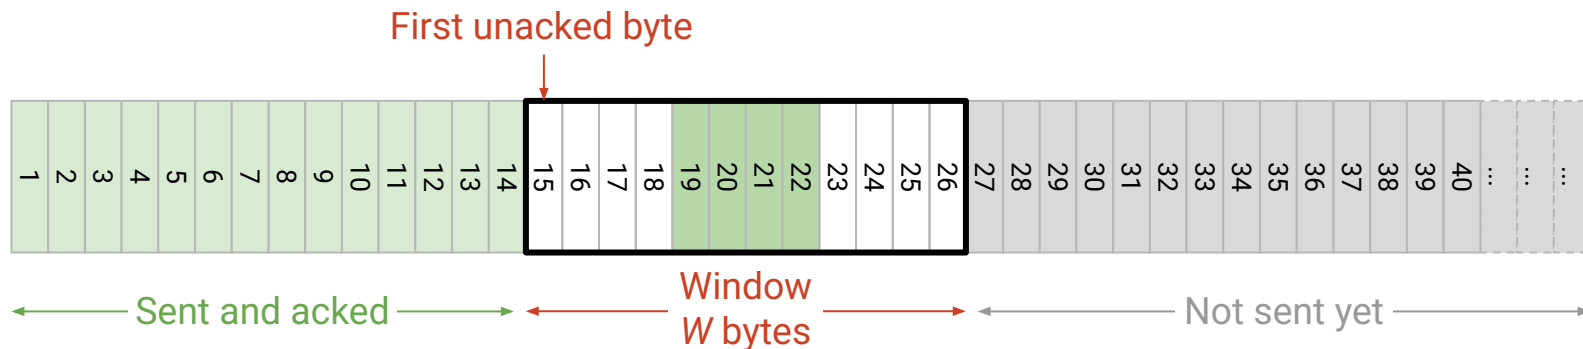
When we measured in packets,  $W$  was the maximum number of packets in flight.

When we measure in bytes,  $W$  is the maximum number of *contiguous* bytes in flight.

- The window is a range of  $W$  contiguous bytes, starting at the first unacked byte.
- Only these  $W$  bytes are allowed to be in flight.

The window slides right if and only if its *leftmost* bytes are acked.

- Acking non-leftmost bytes in the window (e.g. 19–22) does not slide the window.
- The window is determined by the first unacked byte (e.g. still 15).



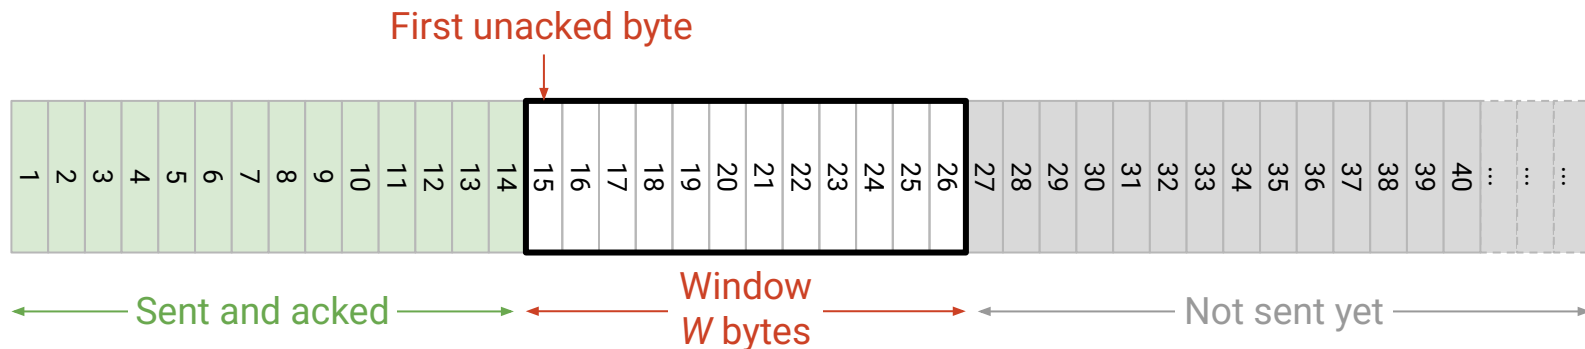
## TCP Sliding Window

TCP uses cumulative acks and a sliding window.

- Cumulative ack: "I have received everything up to (not including) 15."
- Thus, first unacked byte is 15, and the sliding window is  $[15, 15 + W]$ .

$W$  is set as the minimum of two values:

- Advertised window (recipient reports their remaining buffer space).
- Congestion window (sender magically finds value to avoid network overload).




## Detecting Loss

Two ways to detect loss and resend. We resend if either condition is true.

In both cases, we always resend the *first unacked packet* (leftmost part of window).

1. Ack-based: If 3 duplicate acks are received, resend.
2. Timer-based: Keep a single timer. If the timer expires, resend.
  - Different from packet-based TCP, where we kept one timer per packet.
  - Set timer by estimating RTT (e.g. measure times between packets and acks, and take a moving average).

Resend 15–18 if timer expires,  
or if 3 copies of ack(15) received.



# TCP Header

---

Lecture 12, CS 168, Spring 2025

## Implementing TCP

- Byte Notation  
(Segments, Sequence Numbers)
- Maintaining State (Full Duplex,  
Connection Setup and Teardown)
- Sliding Window
- **Header**

What functions does TCP implement?

1. **Demultiplexing** (*ports*)
2. **Reliability** (*checksum, sequence and ack numbers*)
3. **Connection setup and teardown** (*flags*)
4. **Flow control** (*advertised window*)

Source Port (16)			Destination Port (16)		
Sequence Number (32)					
Acknowledgment Number (32)					
Hdr Len (4)	0000	Flags (8)		Advertised Window (16)	
Checksum (8)			Urgent Pointer (8)		
Options (variable-length)					
Payload					

There are 8 flags that can be set in TCP. We care about 4 of them:

- SYN: I'm sending my initial sequence number.
- ACK: I'm acking data (please look at the ack number).
- FIN: I'm done sending data, but will keep receiving data.
- RST: I'm done sending and receiving data.

We won't look at the other four: CWR, ECE, URG, PSH.

Source Port (16)			Destination Port (16)		
Sequence Number (32)					
Acknowledgment Number (32)					
Hdr Len (4)	0000	Flags (8)		Advertised Window (16)	
Checksum (8)			Urgent Pointer (8)		
Options (variable-length)					
Payload					

Remaining fields:

- Header length: Measured in 4-byte words.
  - If no options, this is 5 (header is 20 bytes long).
- 0000: Reserved bits (always set to 0).
- Urgent pointer: Used with the URG flag to indicate urgent data. Won't discuss.
- Options: Extra functionality.
  - Example: SACK uses the options field to implement full-information acks.

Source Port (16)			Destination Port (16)		
Sequence Number (32)					
Acknowledgment Number (32)					
Hdr Len (4)	0000	Flags (8)		Advertised Window (16)	
Checksum (8)			Urgent Pointer (8)		
Options (variable-length)					
Payload					

An elegant (though not perfect) piece of engineering that has stood the test of time.


- Thought experiment: Will TCP continue to be a good solution?

Plenty of evolution in individual pieces:

- Congestion control was added after-the-fact.
- Better acknowledgments, ISN selection, timer estimation, etc.

But the core architectural decisions and abstractions remain:

- Bytestreams, connection-oriented, windows, etc.



Vint Cerf and Bob Kahn have won basically every award ever for their work on TCP.